

Generative AI in Cybersecurity

Module 2B: Retrieval Augmented Generation

Henning Thomsen

htth@ucn.dk

5. May 2025

Pba IT-security @ UCN

Agenda

- Retrieval Augmented Generation (RAG)
 - Document loading
 - Chunking
 - Vector stores
 - Document querying

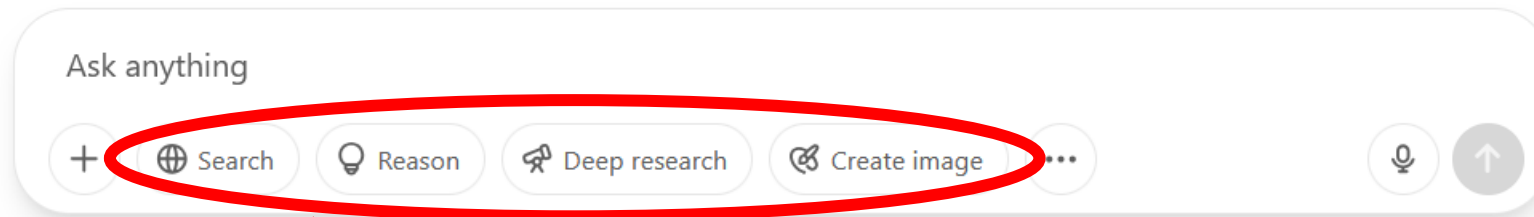
Retrieval Augmented Generation

Augment LLM with (relevant) results from a database

ChatGPT

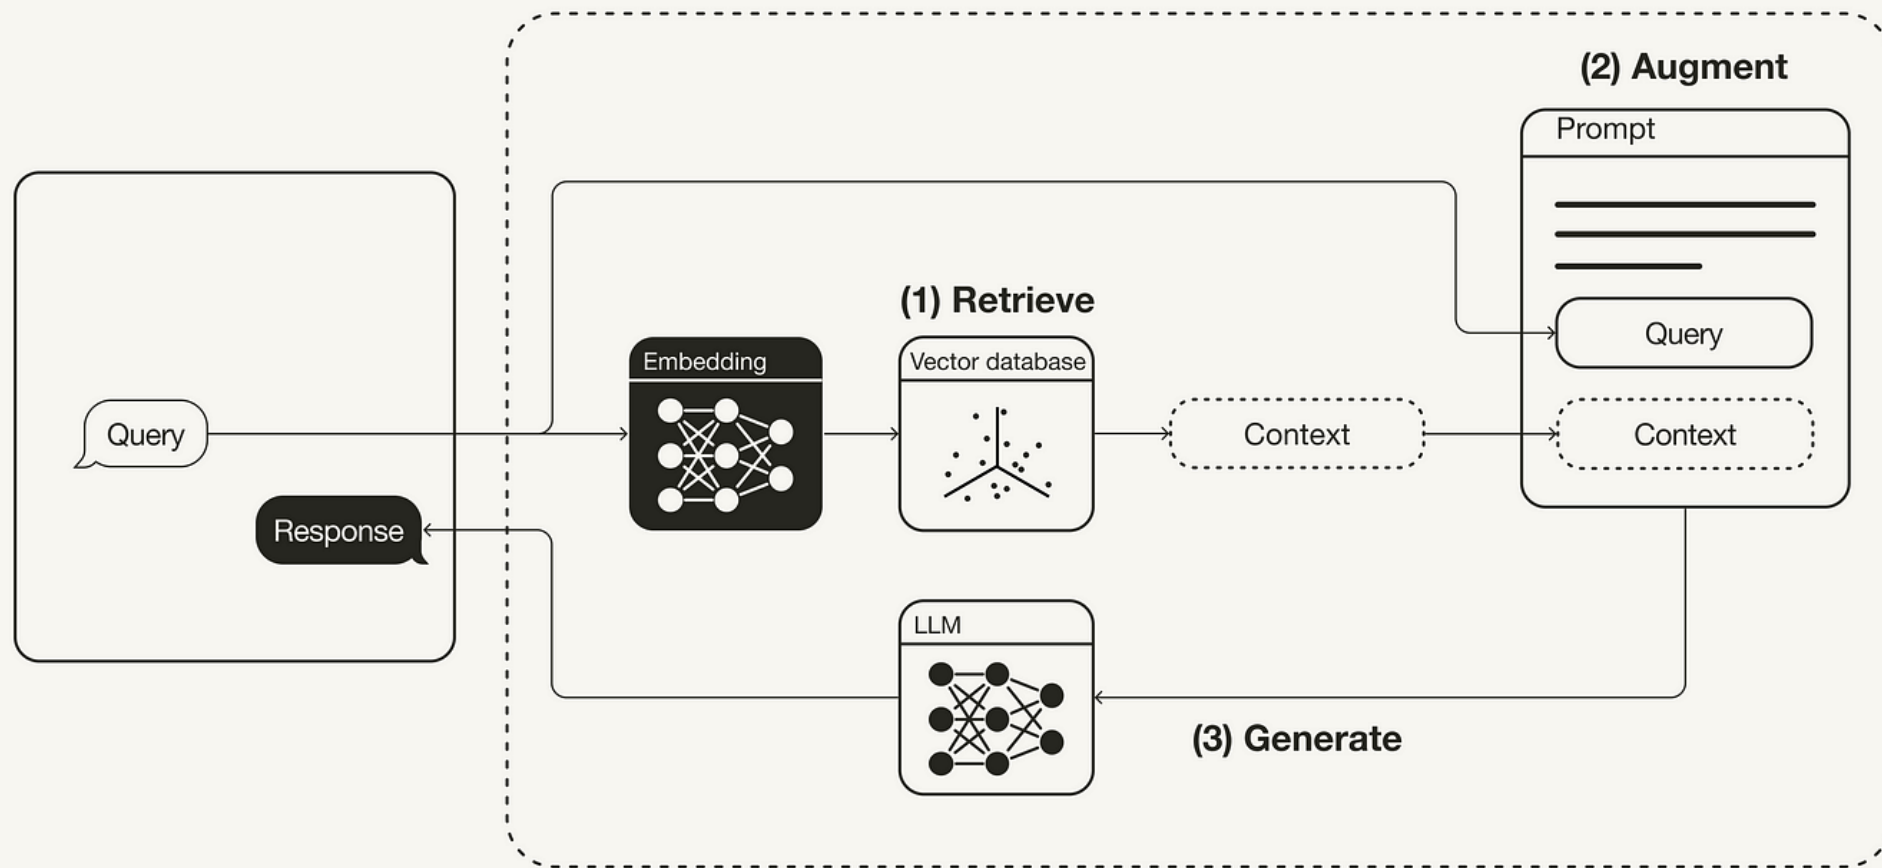


What can I help with?



Extra **tools** added

What is RAG?



Recall spam example

```
prompt_template = ""  
Classify the given email message as either spam or legitimate.
```

Examples are given below:

```
Message: "Hi Alex, just confirming our meeting tomorrow at 10 AM-let me  
know if anything changes."  
Classification: Legitimate
```

```
Message: "Your account has been compromised-click here immediately to  
verify your identity and avoid suspension!"  
Classification: Spam
```

```
Message: {message}  
Classification:  
""
```

Limitations

- Context length related to token cost
- Limited context length (is it sufficient)?
- **Size of context vs. maximum context window**

What is RAG?

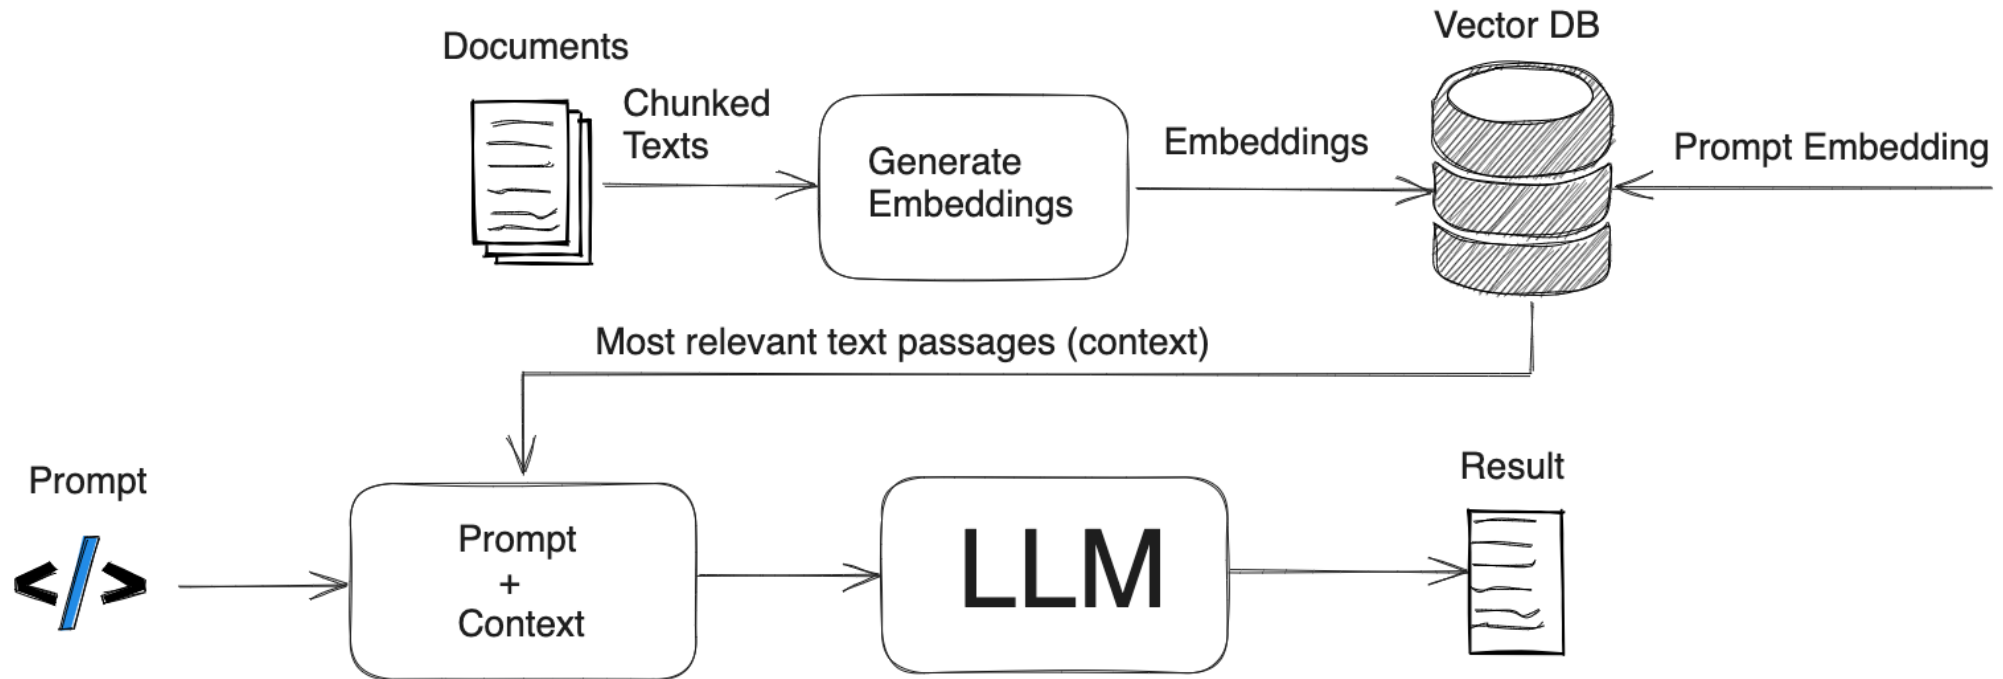


Figure from <https://safjan.com/understanding-retrieval-augmented-generation-rag-empowering-llms/>

Vector embeddings

- LLMs only accept floating point values as input
 - At the most basic level
- Want to convert our context (data) to these vectors
 - Could be: TXT, PDF, DOC, XLS etc.

Vector embeddings

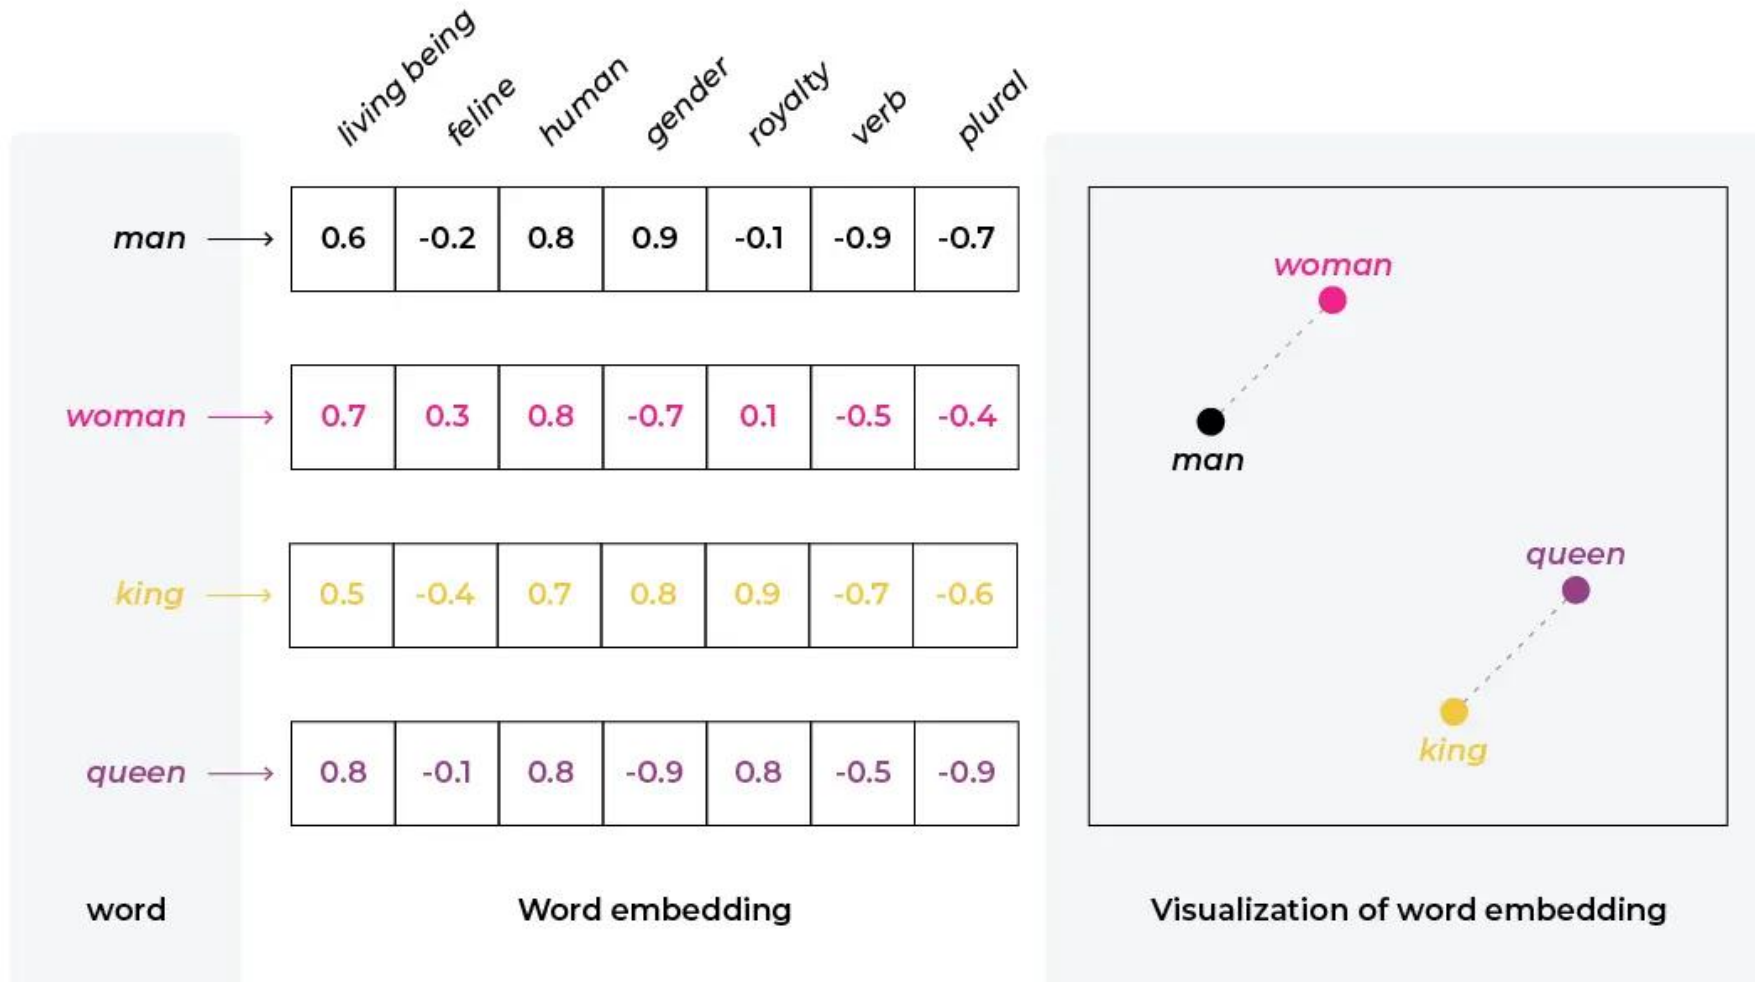



Figure from <https://arize.com/blog-course/embeddings-meaning-examples-and-how-to-compute/>

Indexing a vector database

- Uses an **embedding** to map documents to vectors
 - OpenAIEmbeddings
 - OllamaEmbeddings
- Need to choose among vector stores
 - Chroma
 - FAISS
 - Qdrant
 - ...

Loading documents



```
from langchain_community.document_loaders import PyPDFLoader

file_path = (
    "A-Survey-of-Large-Language-Models.pdf"
)
loader = PyPDFLoader(file_path)
pages = loader.load()
```

Splitting documents



```
from langchain_text_splitters import RecursiveCharacterTextSplitter

text_splitter = RecursiveCharacterTextSplitter(
    chunk_size=1000, chunk_overlap=200, add_start_index=True
)
all_splits = text_splitter.split_documents(pages)
```

Demo: <https://chunkviz.up.railway.app/>


Vector DB ingestion



```
from langchain_openai import OpenAIEmbeddings

from langchain_community.vectorstores import Chroma


vectorstore = Chroma.from_documents(documents=all_splits,
embedding=OpenAIEmbeddings())
```



Querying documents from DB



```
retriever = vectorstore.as_retriever(search_type="similarity",  
search_kwargs={"k": 6})  
  
retrieved_docs = retriever.invoke("What are neural language models?")
```



A simple RAG application using LCEL



```
vectorstore = Chroma.from_documents(documents=all_splits,  
embedding=OpenAIEmbeddings())  
retriever = vectorstore.as_retriever()  
  
chain = (  
    {"context": retriever, "question": RunnablePassthrough()}  
    | prompt  
    | llm  
    | StrOutputParser()  
    )  
  
result = chain.invoke("What are neural language models?")
```